

C++ for High-Frequency Trading Systems

Complete Guide: Theory, Implementation, and Low-Latency Optimization

Chapter 1: Modern C++ Fundamentals for HFT

1.1 Why C++ for High-Frequency Trading?

Performance Requirements:

- **Latency:** Microsecond (μ s) to nanosecond (ns) time scales
- **Throughput:** Millions of operations per second
- **Determinism:** Predictable, consistent execution times
- **Hardware control:** Direct memory management, CPU cache optimization

C++ Advantages:

1. **Zero-cost abstractions:** High-level features without runtime overhead
2. **Manual memory management:** No garbage collection pauses
3. **Inline assembly:** Direct CPU instruction access when needed
4. **Compile-time optimization:** Templates, constexpr eliminate runtime work

1.2 Modern C++ Standards (C++11/14/17/20)

C++11 Key Features:

```
// Move semantics - avoid unnecessary copies
class OrderBook {
    std::vector<Order> orders;
public:
    // Move constructor
    OrderBook(OrderBook&& other) noexcept
        : orders(std::move(other.orders)) {}

    // Move assignment
    OrderBook&& operator=(OrderBook&& other) noexcept {
        orders = std::move(other.orders);
        return *this;
    }
};

// Lambda expressions
auto price_filter = [threshold](const Order& o) {
    return o.price > threshold;
};
```

```

// Smart pointers - automatic memory management
std::unique_ptr<Strategy> strategy = std::make_unique<MomentumStrategy>();
std::shared_ptr<MarketData> data = std::make_shared<MarketData>();

// Range-based for loops
for (const auto& order : orderbook) {
    process(order);
}

```

C++14 Enhancements:

```

// Generic lambdas
auto generic_comparator = [](auto a, auto b) { return a < b; };

// Return type deduction
auto calculate_vwap(const std::vector<Trade> & trades) {
    double sum_price_volume = 0.0;
    double total_volume = 0.0;
    for (const auto& t : trades) {
        sum_price_volume += t.price * t.volume;
        total_volume += t.volume;
    }
    return sum_price_volume / total_volume;
}

```

C++17 Features:

```

// Structured bindings
std::map<std::string, double> positions;
for (const auto& [symbol, quantity] : positions) {
    std::cout << symbol << ":" << quantity << "\n";
}

// std::optional - handle missing values without exceptions
std::optional<Order> find_order(uint64_t order_id) {
    if (auto it = orders.find(order_id); it != orders.end()) {
        return it->second;
    }
    return std::nullopt;
}

// if with initializer
if (auto result = try_execute_order(); result.has_value()) {
    process(result.value());
}

```

C++20 Concepts:

```

// Concepts - compile-time type constraints
template<typename T>;
concept Priceable = requires(T t) {

```

```

    { t.get_price() } -> std::convertible_to<double>;
};

template<Priceable T>;
double calculate_mid(const T& bid, const T& ask) {
    return (bid.get_price() + ask.get_price()) / 2.0;
}

// Ranges
#include <ranges>;
auto expensive_orders = orders
    | std::views::filter([](const auto& o) { return o.price > 1000; })
    | std::views::take(10);

```

Chapter 2: STL Containers for Trading Systems

2.1 Container Selection Guide

Performance Characteristics:

Container	Insert	Lookup	Delete	Use Case
std::vector	O(1) amortized end	O(n)	O(n)	Time-series data, tick storage
std::deque	O(1) both ends	O(n)	O(1) ends	Ring buffers, sliding windows
std::list	O(1) anywhere	O(n)	O(1)	Order lists (price levels)
std::map	O(log n)	O(log n)	O(log n)	Price levels (Red-Black tree)
std::unordered_map	O(1) average	O(1) average	O(1) average	Symbol lookups, order IDs
std::set	O(log n)	O(log n)	O(log n)	Unique sorted elements

2.2 Order Book Implementation

Limit Order Book Structure:

```

#include <map>;
#include <list>;
#include <unordered_map>

struct Order {
    uint64_t order_id;
    uint64_t timestamp;
    char side; // 'B' for buy, 'S' for sell
    double price;
    uint32_t quantity;

    Order(uint64_t id, uint64_t ts, char s, double p, uint32_t q)
        : order_id(id), timestamp(ts), side(s), price(p), quantity(q) {}

};


```

```

class OrderBook {
private:
    // Price level -> list of orders (FIFO queue)
    using OrderList = std::list<Order>;
    using PriceLevel = std::pair<uint32_t, OrderList>; // (total_quantity, orders)

    // Buy side: highest price first (descending)
    std::map<double, PriceLevel, std::greater<double>> bids;

    // Sell side: lowest price first (ascending)
    std::map<double, PriceLevel, std::less<double>> asks;

    // Fast order lookup by ID
    std::unordered_map<uint64_t, std::pair<double, OrderList::iterator>> order_map;
    uint64_t total_qty = 0;

public:
    // Add order: O(log n) for price level + O(1) for order insertion
    void add_order(const Order& order) {
        auto& levels = (order.side == 'B') ? bids : asks;
        auto& [total_qty, order_list] = levels[order.price];

        order_list.push_back(order);
        total_qty += order.quantity;

        // Store iterator for O(1) deletion
        order_map[order.order_id] = {order.price, std::prev(order_list.end())};
    }

    // Cancel order: O(log n) + O(1)
    bool cancel_order(uint64_t order_id) {
        auto it = order_map.find(order_id);
        if (it == order_map.end()) return false;

        auto [price, order_it] = it->second;
        char side = order_it->side;

        auto& levels = (side == 'B') ? bids : asks;
        auto level_it = levels.find(price);

        auto& [total_qty, order_list] = level_it->second;
        total_qty -= order_it->quantity;
        order_list.erase(order_it);

        // Remove empty price level
        if (order_list.empty()) {
            levels.erase(level_it);
        }

        order_map.erase(it);
        return true;
    }

    // Get best bid/ask: O(1)
    std::optional<double> best_bid() const {
        if (bids.empty()) return std::nullopt;
        return bids.begin()->first;
    }
}

```

```

    }

    std::optional<double> best_ask() const {
        if (asks.empty()) return std::nullopt;
        return asks.begin()->first;
    }

    // Match incoming order
    std::vector<std::pair<uint64_t, uint32_t>> match_order(Order& order, incoming);
    std::vector<std::pair<uint64_t, uint32_t>> fills;

    auto& levels = (incoming.side == 'B') ? asks : bids;

    while (!levels.empty() && incoming.quantity > 0) {
        auto& [price, level] = *levels.begin();
        auto& [total_qty, order_list] = level;

        // Check price crossing
        bool crosses = (incoming.side == 'B' && incoming.price >= price) ||
                       (incoming.side == 'S' && incoming.price <= price);
        if (!crosses) break;

        // Match with orders at this price level (FIFO)
        while (!order_list.empty() && incoming.quantity > 0) {
            Order& resting = order_list.front();
            uint32_t fill_qty = std::min(incoming.quantity, resting.quantity);

            fills.push_back({resting.order_id, fill_qty});

            incoming.quantity -= fill_qty;
            resting.quantity -= fill_qty;
            total_qty -= fill_qty;

            if (resting.quantity == 0) {
                order_map.erase(resting.order_id);
                order_list.pop_front();
            }
        }
    }

    // Remove empty price level
    if (order_list.empty()) {
        levels.erase(levels.begin());
    }
}

return fills;
}
};


```

Chapter 3: Multithreading and Concurrency

3.1 Thread Basics

Creating Threads:

```
#include <thread>
#include <iostream>

void market_data_handler() {
    while (running) {
        // Process incoming market data
        process_tick();
    }
}

void order_execution_handler() {
    while (running) {
        // Execute pending orders
        execute_orders();
    }
}

int main() {
    std::thread md_thread(market_data_handler);
    std::thread exec_thread(order_execution_handler);

    // Wait for threads to complete
    md_thread.join();
    exec_thread.join();
}
```

3.2 Thread Synchronization

Mutex and Lock Guards:

```
#include <mutex>

class ThreadSafeOrderBook {
private:
    OrderBook orderbook;
    mutable std::mutex mtx;

public:
    void add_order(const Order& order) {
        std::lock_guard<std::mutex> lock(mtx); // RAI-style locking
        orderbook.add_order(order);
    } // Automatically unlocks when lock_guard destructs

    std::optional<double> best_bid() const {
        std::lock_guard<std::mutex> lock(mtx);
        return orderbook.best_bid();
    }
}
```

```
    }
};
```

Condition Variables (Producer-Consumer):

```
#include <condition_variable>;
#include <queue>

class ThreadSafeQueue {
private:
    std::queue<Order> queue;
    mutable std::mutex mtx;
    std::condition_variable cv;

public:
    void push(Order order) {
        {
            std::lock_guard<std::mutex> lock(mtx);
            queue.push(std::move(order));
        }
        cv.notify_one(); // Wake up waiting consumer
    }

    Order pop() {
        std::unique_lock<std::mutex> lock(mtx);
        cv.wait(lock, [this] { return !queue.empty(); }); // Wait for data

        Order order = std::move(queue.front());
        queue.pop();
        return order;
    }
};
```

3.3 Lock-Free Programming with std::atomic

Atomic Operations:

```
#include <atomic>

class LockFreeCounter {
private:
    std::atomic<uint64_t> counter{0};

public:
    uint64_t increment() {
        return counter.fetch_add(1, std::memory_order_relaxed);
    }

    uint64_t get() const {
        return counter.load(std::memory_order_acquire);
    }
};
```

Lock-Free SPSC (Single Producer Single Consumer) Queue:

```
template<typename T, size_t Size>;
class SPSCQueue {
private:
    std::array<T, Size> buffer;
    alignas(64) std::atomic<size_t> write_idx{0}; // Separate cache lines
    alignas(64) std::atomic<size_t> read_idx{0};

public:
    bool push(const T& item) {
        size_t current_write = write_idx.load(std::memory_order_relaxed);
        size_t next_write = (current_write + 1) % Size;

        // Check if queue is full
        if (next_write == read_idx.load(std::memory_order_acquire)) {
            return false; // Queue full
        }

        buffer[current_write] = item;
        write_idx.store(next_write, std::memory_order_release);
        return true;
    }

    bool pop(T& item) {
        size_t current_read = read_idx.load(std::memory_order_relaxed);

        // Check if queue is empty
        if (current_read == write_idx.load(std::memory_order_acquire)) {
            return false; // Queue empty
        }

        item = buffer[current_read];
        read_idx.store((current_read + 1) % Size, std::memory_order_release);
        return true;
    }
};
```

Memory Ordering Explanation:

- `memory_order_relaxed`: No synchronization, only atomicity
- `memory_order_acquire`: Reads happen-before subsequent operations
- `memory_order_release`: Writes happen-after previous operations
- `memory_order_acq_rel`: Both acquire and release
- `memory_order_seq_cst`: Sequential consistency (strongest, default)

Chapter 4: Memory Management and Optimization

4.1 Custom Allocators

Object Pool Allocator:

```
template<typename T, size_t PoolSize = 1024>;
class ObjectPool {
private:
    union Slot {
        T object;
        Slot* next;
    };

    std::array<Slot, PoolSize> pool;
    Slot* free_list;

public:
    ObjectPool() {
        // Initialize free list
        free_list = &pool[0];
        for (size_t i = 0; i < PoolSize - 1; ++i) {
            pool[i].next = &pool[i + 1];
        }
        pool[PoolSize - 1].next = nullptr;
    }

    template<typename... Args>;
    T* allocate(Args&... args) {
        if (!free_list) return nullptr; // Pool exhausted

        Slot* slot = free_list;
        free_list = free_list->next;

        // Placement new
        return new (&slot->object) T(std::forward<Args>(args)...);
    }

    void deallocate(T* ptr) {
        ptr->~T(); // Explicit destructor call

        Slot* slot = reinterpret_cast<Slot*>(ptr);
        slot->next = free_list;
        free_list = slot;
    }
};

// Usage
ObjectPool<Order> order_pool;
Order* order = order_pool.allocate(12345, timestamp, 'B', 100.5, 1000);
// ... use order ...
order_pool.deallocate(order);
```

4.2 Cache Optimization

False Sharing Prevention:

```
// BAD: False sharing (both on same cache line)
struct BadCounter {
    std::atomic<uint64_t> producer_count;
    std::atomic<uint64_t> consumer_count;
};

// GOOD: Separate cache lines (64 bytes)
struct alignas(64) GoodCounter {
    std::atomic<uint64_t> producer_count;
    char padding[64 - sizeof(std::atomic<uint64_t>)];
};

struct alignas(64) ConsumerCounter {
    std::atomic<uint64_t> consumer_count;
};
```

Cache-Friendly Data Layout (SoA vs AoS):

```
// Array of Structures (AoS) - poor cache utilization
struct Trade {
    double price;
    uint32_t volume;
    uint64_t timestamp;
};
std::vector<Trade> trades; // Bad for iterating over just prices

// Structure of Arrays (SoA) - better cache utilization
struct TradeData {
    std::vector<double> prices;
    std::vector<uint32_t> volumes;
    std::vector<uint64_t> timestamps;
};
TradeData trades; // Good: prices contiguous in memory
```

4.3 Memory Alignment

```
// Align data for SIMD operations
alignas(32) float prices[8]; // Aligned for AVX (256-bit)

// Check alignment
static_assert(alignof(prices) == 32);

// Aligned allocation
void* aligned_alloc_wrapper(size_t size, size_t alignment) {
    void* ptr = nullptr;
    if (posix_memalign(&ptr, alignment, size) != 0) {
        throw std::bad_alloc();
    }
}
```

```
    return ptr;  
}
```

Chapter 5: SIMD and Vectorization

5.1 Introduction to SIMD

SIMD = Single Instruction Multiple Data

- Process multiple data elements with one CPU instruction
- Modern CPUs: SSE (128-bit), AVX/AVX2 (256-bit), AVX-512 (512-bit)

5.2 AVX2 Intrinsics

Vectorized Dot Product:

```
#include <immintrin.h> // AVX/AVX2 intrinsics  
  
// Dot product of two arrays (AVX2: 8 floats at once)  
float dot_product_avx2(const float* a, const float* b, size_t n) {  
    __m256 sum = _mm256_setzero_ps(); // Zero vector  
  
    size_t i = 0;  
    for (; i + 7 < n; i += 8) {  
        __m256 va = _mm256_loadu_ps(&a[i]); // Load 8 floats  
        __m256 vb = _mm256_loadu_ps(&b[i]);  
        __m256 prod = _mm256_mul_ps(va, vb); // Element-wise multiply  
        sum = _mm256_add_ps(sum, prod); // Accumulate  
    }  
  
    // Horizontal sum: sum all 8 elements in sum vector  
    __m128 sum_low = _mm256_castps256_ps128(sum);  
    __m128 sum_high = _mm256_extractf128_ps(sum, 1);  
    __m128 sum128 = _mm_add_ps(sum_low, sum_high);  
  
    sum128 = _mm_hadd_ps(sum128, sum128);  
    sum128 = _mm_hadd_ps(sum128, sum128);  
  
    float result = _mm_cvtsd_f32(sum128);  
  
    // Handle remaining elements (scalar)  
    for (; i < n; ++i) {  
        result += a[i] * b[i];  
    }  
  
    return result;  
}
```

VWAP Calculation with AVX2:

```

double calculate_vwap_avx2(const double* prices, const double* volumes, size_t n) {
    __m256d sum_pv = _mm256_setzero_pd(); // 4 doubles
    __m256d sum_vol = _mm256_setzero_pd();

    size_t i = 0;
    for (; i + 3 < n; i += 4) {
        __m256d p = _mm256_loadu_pd(&prices[i]);
        __m256d v = _mm256_loadu_pd(&volumes[i]);

        __m256d pv = _mm256_mul_pd(p, v);
        sum_pv = _mm256_add_pd(sum_pv, pv);
        sum_vol = _mm256_add_pd(sum_vol, v);
    }

    // Horizontal sum
    double pv_array[4], vol_array[4];
    _mm256_storeu_pd(pv_array, sum_pv);
    _mm256_storeu_pd(vol_array, sum_vol);

    double total_pv = pv_array[0] + pv_array[1] + pv_array[2] + pv_array[3];
    double total_vol = vol_array[0] + vol_array[1] + vol_array[2] + vol_array[3];

    // Scalar remainder
    for (; i < n; ++i) {
        total_pv += prices[i] * volumes[i];
        total_vol += volumes[i];
    }

    return total_pv / total_vol;
}

```

Chapter 6: Template Metaprogramming

6.1 Templates for Generic Code

Function Templates:

```

template<typename T>;
T max_price(const T& a, const T& b) {
    return (a > b) ? a : b;
}

// Usage
double max_d = max_price(100.5, 101.2);
int max_i = max_price(100, 200);

```

Class Templates:

```

template<typename T, size_t Capacity>;
class CircularBuffer {
private:

```

```

    std::array<T, Capacity> buffer;
    size_t head = 0;
    size_t tail = 0;
    size_t count = 0;

public:
    bool push(const T& item) {
        if (count == Capacity) return false;
        buffer[tail] = item;
        tail = (tail + 1) % Capacity;
        ++count;
        return true;
    }

    bool pop(T& item) {
        if (count == 0) return false;
        item = buffer[head];
        head = (head + 1) % Capacity;
        --count;
        return true;
    }

    constexpr size_t capacity() const { return Capacity; }
};


```

6.2 CRTP (Curiously Recurring Template Pattern)

Static Polymorphism (Zero Runtime Overhead):

```

template<typename Derived>;
class Strategy {
public:
    void execute() {
        static_cast<Derived*>(this)->execute_impl();
    }

    double calculate_signal() {
        return static_cast<Derived*>(this)->calculate_signal_impl();
    }
};

class MomentumStrategy : public Strategy<MomentumStrategy> {
public:
    void execute_impl() {
        // Momentum-specific logic
    }

    double calculate_signal_impl() {
        return momentum_indicator;
    }
};

class MeanReversionStrategy : public Strategy<MeanReversionStrategy> {
public:
    void execute_impl() {

```

```

        // Mean reversion logic
    }

    double calculate_signal_impl() {
        return z_score;
    }
};

// Generic function works with any Strategy
template<typename T>;
void run_strategy(Strategy<T> strategy) {
    double signal = strategy.calculate_signal();
    if (signal > threshold) {
        strategy.execute();
    }
}

```

6.3 Compile-Time Computation

constexpr Functions:

```

constexpr uint64_t factorial(int n) {
    return (n <= 1) ? 1 : n * factorial(n - 1);
}

constexpr uint64_t combinations(int n, int k) {
    return factorial(n) / (factorial(k) * factorial(n - k));
}

// Computed at compile time!
constexpr uint64_t result = combinations(10, 3); // No runtime cost

```

Chapter 7: Low-Latency Techniques

7.1 Latency Sources

1. **Algorithmic complexity:** $O(n)$ vs $O(1)$
2. **Cache misses:** L1 (~1ns) vs L2 (~3ns) vs L3 (~12ns) vs RAM (~100ns)
3. **Branch misprediction:** ~15-20 cycles penalty
4. **Context switches:** Microseconds
5. **System calls:** Avoid kernel transitions
6. **Network:** Physical distance, kernel stack

7.2 Branch Prediction Hints

```
// Likely/unlikely macros (GCC/Clang)
#define LIKELY(x) __builtin_expect(!!(x), 1)
#define UNLIKELY(x) __builtin_expect(!!(x), 0)

if (LIKELY(order.quantity > 0)) {
    // Hot path - processor prefetches this branch
    execute_order(order);
} else {
    // Cold path
    handle_error();
}
```

7.3 CPU Affinity (Pin Thread to Core)

```
#include <pthread.h>
#include <sched.h>

void set_thread_affinity(int core_id) {
    cpu_set_t cpuset;
    CPU_ZERO(&cpuset);
    CPU_SET(core_id, &cpuset);

    pthread_t current_thread = pthread_self();
    pthread_setaffinity_np(current_thread, sizeof(cpu_set_t), &cpuset);
}

// Pin critical thread to core 0
set_thread_affinity(0);
```

7.4 Timestamping with RDTSC

```
#include <x86intrin.h>

inline uint64_t rdtsc() {
    return __rdtsc(); // Read CPU Time Stamp Counter
}

// Measure latency
uint64_t start = rdtsc();
process_order(order);
uint64_t end = rdtsc();
uint64_t cycles = end - start;

// Convert to nanoseconds (assuming 3 GHz CPU)
double ns = cycles / 3.0;
```

Chapter 8: Networking for HFT

8.1 TCP Socket Programming

Basic TCP Server:

```
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>
#include <unistd.h>

int create_tcp_server(int port) {
    int server_fd = socket(AF_INET, SOCK_STREAM, 0);

    int opt = 1;
    setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR, &opt, sizeof(opt));

    sockaddr_in address;
    address.sin_family = AF_INET;
    address.sin_addr.s_addr = INADDR_ANY;
    address.sin_port = htons(port);

    bind(server_fd, (sockaddr*)&address, sizeof(address));
    listen(server_fd, 10);

    return server_fd;
}
```

8.2 UDP Multicast (Market Data Feeds)

Multicast Receiver:

```
int create_multicast_receiver(const char* group, int port) {
    int sock = socket(AF_INET, SOCK_DGRAM, 0);

    // Allow multiple sockets to bind to same port
    int reuse = 1;
    setsockopt(sock, SOL_SOCKET, SO_REUSEADDR, &reuse, sizeof(reuse));

    sockaddr_in local_addr;
    memset(&local_addr, 0, sizeof(local_addr));
    local_addr.sin_family = AF_INET;
    local_addr.sin_addr.s_addr = INADDR_ANY;
    local_addr.sin_port = htons(port);

    bind(sock, (sockaddr*)&local_addr, sizeof(local_addr));

    // Join multicast group
    ip_mreq mreq;
    mreq.imr_multicast.s_addr = inet_addr(group);
    mreq.imr_interface.s_addr = INADDR_ANY;
    setsockopt(sock, IPPROTO_IP, IP_ADD_MEMBERSHIP, &mreq, sizeof(mreq));
```

```
    return sock;  
}
```

8.3 Kernel Bypass with DPDK (Conceptual)

DPDK = Data Plane Development Kit

- Bypass kernel network stack
- Poll-mode drivers (no interrupts)
- User-space packet processing
- Achieve <1 μ s latency

Key Concepts:

- **Huge pages:** Reduce TLB misses
- **CPU isolation:** Dedicated cores for packet processing
- **Zero-copy:** DMA directly to user space
- **Batch processing:** Process packets in bursts

Chapter 9: Project Implementations

Project 1: Lock-Free Order Book

Difficulty: Advanced

Technologies: std::atomic, Red-Black tree, SPSC queue

Metrics: 450K+ operations/second, <100ns latency

Project 2: Market Data Handler

Difficulty: Advanced

Technologies: UDP multicast, FIX protocol, multithreading

Metrics: 10M+ messages/second, <5 μ s tick-to-trade

Project 3: Matching Engine Simulator

Difficulty: Intermediate-Advanced

Features: FIFO matching, price-time priority, trade reporting

Project 4: Custom Memory Pool

Difficulty: Intermediate

Features: Object pooling, NUMA-aware allocation, zero fragmentation

Project 5: SIMD Matrix Operations Library

Difficulty: Advanced

Features: AVX2 vectorization, cache blocking, 10x+ speedup

Chapter 10: Interview Preparation

10.1 Common C++ Interview Questions

Question 1: Implement std::unique_ptr

```
template<typename T>
class UniquePtr {
private:
    T* ptr;

public:
    explicit UniquePtr(T* p = nullptr) : ptr(p) {}

    ~UniquePtr() { delete ptr; }

    // Delete copy constructor and assignment
    UniquePtr(const UniquePtr&) = delete;
    UniquePtr& operator=(const UniquePtr&) = delete;

    // Move constructor
    UniquePtr(UniquePtr&& other) noexcept : ptr(other.ptr) {
        other.ptr = nullptr;
    }

    // Move assignment
    UniquePtr& operator=(UniquePtr&& other) noexcept {
        if (this != &other) {
            delete ptr;
            ptr = other.ptr;
            other.ptr = nullptr;
        }
        return *this;
    }

    T& operator*() const { return *ptr; }
    T* operator->() const { return ptr; }
    T* get() const { return ptr; }

    T* release() {
        T* old_ptr = ptr;
        ptr = nullptr;
        return old_ptr;
    }

    void reset(T* p = nullptr) {
        delete ptr;
        ptr = p;
    }
}
```

```
    }  
};
```

Question 2: Explain Virtual Function Mechanism

- Virtual table (vtable): Array of function pointers
- Virtual pointer (vptr): Hidden pointer in each object
- Runtime polymorphism cost: One extra indirection (~3-5 ns)
- Can prevent inlining and optimization

Question 3: What is False Sharing?

- Multiple threads modify different variables on same cache line
- Cache line (64 bytes) invalidated on each write
- Solution: Align variables to separate cache lines (`alignas(64)`)

10.2 System Design Questions

Question: Design a High-Performance Order Management System

Components:

1. **Order Ingestion:** Multithreaded receivers, lock-free queues
2. **Risk Checks:** Pre-trade validation (limits, exposure)
3. **Order Routing:** Smart order routing to exchanges
4. **Execution:** FIX protocol, TCP/UDP connections
5. **Position Tracking:** Real-time P&L calculation
6. **Logging:** Async logging to avoid blocking critical path

Performance Targets:

- Order-to-wire: <10 µs
- Throughput: 100K+ orders/second
- Uptime: 99.999% (5 nines)

Further Reading

Books:

- Scott Meyers: *Effective Modern C++*
- Anthony Williams: *C++ Concurrency in Action*
- Fedor Pikus: *The Art of Writing Efficient Programs*
- Agner Fog: *Optimizing Software in C++* (free PDF)

Online Resources:

- [CPPReference.com](#) - C++ standard library reference
- Compiler Explorer ([godbolt.org](#)) - Assembly output analysis
- Intel Intrinsics Guide - SIMD instruction reference
- CppCon YouTube - Conference talks on performance

Open Source Projects:

- Folly (Facebook) - High-performance C++ library
- Abseil (Google) - C++ common libraries
- DPDK - Data Plane Development Kit
- libfixpp - FIX protocol parser

This textbook provides comprehensive coverage of C++ for HFT systems, from fundamentals to advanced low-latency techniques, with complete working implementations and interview preparation materials.